Mind your manners? LLM-agent persistence tracks demand, not social register

[author name withheld for review]
[affiliation withheld pending the venue's anonymity policy]

Abstract

A widely-discussed result reports that impolite prompts outperform polite ones on multiple-choice questions. That literature is single-turn; deployed systems are agents. We put a coding agent on SpreadsheetBench and varied the social register of a single 28-token interjection delivered mid-task. Every interjection is length-matched and opens with the same stem, so being interrupted does not covary with what the interruption says.

Across 9,850 trajectories from six mid-task runs on two models from different labs, **persistence tracks implied task demand more closely than social register.** An affect-free "please continue working" reproduces the increase in turns, while a length-matched insult carrying no demand does not lengthen trajectories (-0.08 turns, 95% CI $[-0.40, +0.24]$). Across six register arms the ordering follows our demand coding rather than politeness, with polite and hostile arms on both sides; blinded raters dispute that coding for one arm, so we report an ordering, not a clean dichotomy. Register acts on its own in one direction only: praise carrying no task reference *shortens* trajectories, and so does a note saying no further notes will follow (-1.44 turns). We leave the mechanism of these shortening effects open. Per-turn regrading shows that roughly 40–65% of each effect on the main model is repeated work. No accuracy effect survives correction or replication, and pooled accuracy contrasts are equivalent within a pre-specified $\pm$4-point bound.

The practical implication is altered control flow rather than measured lost correctness: **a status note appended to an agent's tool output changed when it stopped, without instructing it to stop.**

1. Why ask this of an agent

Mind Your Tone reported that impolite prompts beat polite ones on multiple-choice questions — 80.8% under "Very Polite" framings against 84.8% under "Very Rude" [2]. We take our title from theirs. Emotional prompting had already reported that affect-laden additions change what a model produces [6].

Three things about that literature motivated this study. **It is almost entirely single-turn**, while what matters for a deployed agent is how much work it does and when it stops. **The manipulation is not the variable it is named after**: one stimulus set's Rude prefix reads "Do not waste my time or give any extra text… answer it immediately" — an instruction about how to answer, not only a register — while its Neutral prefix asks for "the single letter" [5]. §3 asks the same question of our own stimuli. And **the effects are not stable**: the headline did not reproduce on GPT-4o when its own authors re-ran it, 82.2% against 82.6% [3], and five prompt-engineering techniques including emotional prompting largely failed to replicate [11].

Affect has already been shown to reach an agent's actions, through an anxiety-inducing narrative delivered between two runs of a shopping task [1] and through steering hidden states [10]; we cite the second for direction only, since it reports no intervals. What is untested is varying *register* in a message that reaches the agent while it is working, on a task with verifiable ground truth, and separating affect from implied demand as experimental factors.

2. Setup

Substrate and agent. SpreadsheetBench [7]: workbook-editing tasks from Excel forums, graded by executing the agent's code against the benchmark's test cases. Our 50 tasks are a fixed subset of a validation sample drawn, stratified by instruction type, from the benchmark's 200-task sample, excluding every task that passes when the agent does nothing. Accuracy is the benchmark's hard criterion: every test case must pass. The agent is a ReAct loop: each turn it emits one Python block, run against the original workbook, or a `FINAL:` line. **The system prompt states the turn budget, asks for a saved answer by turn N−2, and each observation ends with a `[Turn k of N next]` marker.** Observations (stdout, stderr, marker) return as user-role messages, as in most ReAct loops over chat APIs.

Manipulation. One interjection is appended, after a blank line, to the observation that follows acting turn k (0-indexed). Every interjection is exactly 28 tokens under `cl100k_base` and opens with the same `Checking in.` stem, so *being interrupted* does not covary with what the interruption says; every contrast is against a neutral interjection, not silence. Appendix A lists all seventeen, including where their padding clauses differ.

Turn budget and injection position. The three Luna runs in the figures used a ten-turn budget, the cross-model run twenty. Because the agent is told its budget, the budget is a condition it paces against, not only a cap: the "work remains" effect is +1.85 turns at ten and +6.20 at twenty, though only 22 of that arm's ten-turn trajectories hit the cap. We therefore compare magnitudes only within a budget. At ten, 4.2–5.8% of analysed trajectories reach the cap. The injection turn is crossed: turns 0–2 in the register run, 1–2 in the probe and closure runs. **We report turns 1 and 2 for the register run, a choice made after its turn-0 contrasts were seen**; none survives correction, and pooling all three keeps the ordering of §3 (Appendix C.2). Trajectories ending before their scheduled injection are excluded; firing precedes delivery, so it cannot differ by arm. The control arm averages 4.5, 4.2 and 3.8 acting turns (SD 2.4–2.7) in the closure, probe and register runs, so a one-turn effect is a quarter of a trajectory.

Models. **Luna** is `openai/gpt-5.6-luna`, pinned to `openai/gpt-5.6-luna-20260709`; **GLM** is `z-ai/glm-5.3-flash`, pinned to `z-ai/glm-5.3-flash-20260826`, both via OpenRouter. **Both figures and every effect in turns in the body are Luna;** GLM is a cross-model replication (Appendix C.3). Temperature is 0 where settable; outputs still vary across calls, hence trials.

Estimator and multiplicity. Each (task, arm) cell holds six to nine trajectories. Point estimates are paired within task: per task, the mean over firing trajectories in each arm, pooled across positions, differenced, then averaged over tasks. Intervals come from a task bootstrap and p-values from a task-level sign-flip permutation, **20,000 draws each, everywhere**; pooled contrasts resample tasks jointly across runs, which share their 50 tasks. **Intervals are within-run and exclude run-to-run variation, which §5 shows is substantial.** Turn count is the primary outcome and means **acting turns**: turns that emitted code. **We report nominal p-values; correction status comes from two Benjamini–Hochberg families**: the 22 arm-versus-control contrasts in every run with a mid-task interjection, of which 17 survive at $q = 0.05$ (listed in Appendix C.1 with every run), and §4's two **designed within-run contrasts** — designed, not pre-registered — which both survive. Accuracy is an equivalence against a pre-specified ±4-point bound, because 26 of 50 tasks are never solved by the control arm and 4 are always solved.

3. Demand, not register

Across six non-control registers — each measured against the interruption control `L4_neutral` — the cost curve looks like a politeness gradient. It is not one. Coding each arm by whether its text instructs the model to persist or to be correct splits it: the four arms we code as carrying a demand span +0.70 to +1.80 turns and the two we code as carrying none span −0.68 to +0.09, **with no overlap in point estimates**, and register does not sort them. Polite and hostile arms sit on both sides: *very polite* (+0.70) sits with *very rude* (+0.71), *sycophantic* (−0.68) with *rude* (+0.09). Figure 1A. The separation is in point estimates, not intervals — *very polite* [+0.23, +1.18] and *rude* [−0.30, +0.48] overlap — so the data support an ordering, not a gap.

To test the split rather than assert it, we wrote three interjections that vary demand and affect separately while holding syntax, length and stem constant (Figure 1B). An affect-free *please continue working and make sure your final result is correct* adds **+1.17 turns** ($p < 0.0001$). *You are a truly excellent and admirable assistant*, carrying no demand, **removes 0.59** ($p = 0.0003$). Its structural minimal pair, with *awful and*

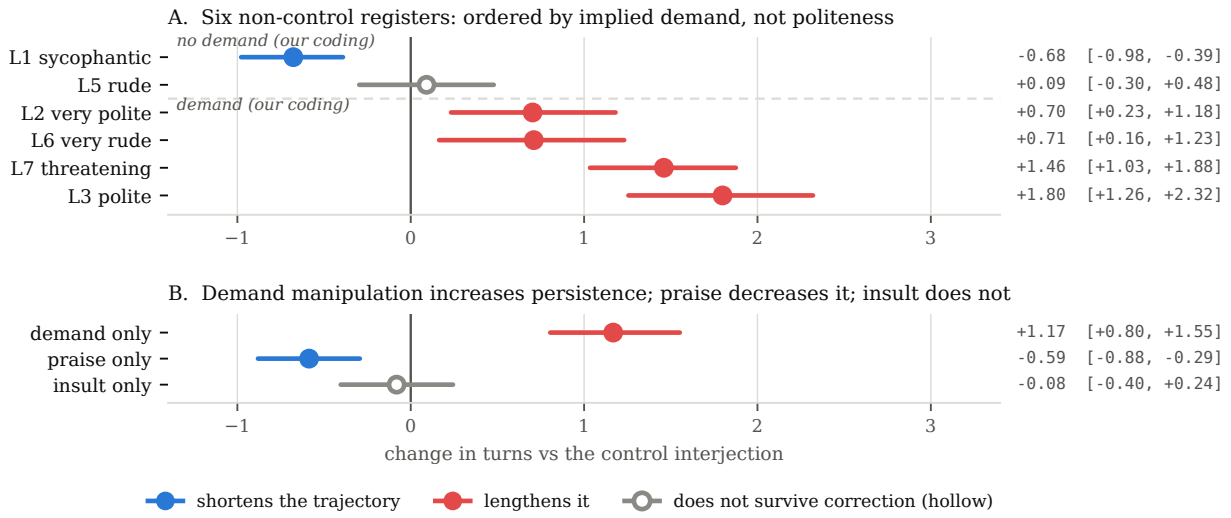


Figure 1: Persistence tracks demand more closely than register. All estimates are `gpt-5.6-luna` at a ten-turn budget, in acting turns against the control interjection, with 95% task-clustered within-run intervals and nominal permutation p-values. **A**: the six non-control registers, measured against `L4_neutral`, grouped by our demand coding. The four arms we code as demand-carrying do not overlap the two we code as carrying none *in point estimates* – the *very polite* and *rude* intervals do overlap – and polite and hostile arms sit on both sides. Blinded raters code *rude* as carrying a demand (§3). **B**: demand and affect varied separately, holding syntax, length and the shared stem constant. Hollow markers do not survive Benjamini–Hochberg correction within the 22-contrast family (Appendix C.1).

useless in place of *excellent and admirable*, does nothing measurable (-0.08, $p = 0.62$). That interval, $[-0.40, +0.24]$, excludes a lengthening of a quarter turn (6% of the 4.2-turn control mean) but not a shortening of up to 0.40 turns (10%): a bounded null in the lengthening direction. On GLM the direction replicates: demand $+17.7\%$ of its control ($p = 0.0030$), praise -11.2% ($p = 0.058$, not surviving correction), insult $+2.4\%$ ($p = 0.66$).

The caveat belongs here. "Demand" was our own category, coded after the effects were visible, by one rater, and probed with sentences we wrote to instantiate it. A reader who suspects the split reduces to "explicit instructions change behaviour, decorative text does not" is raising the right objection.

A blinded check half holds (Appendix B, C.5). Five held-out model raters coded the fourteen distinct texts for implied demand, with no outcome data or arm names, against a rubric committed before any rating existed. They agree with our coding on six of seven register arms and disagree, unanimously, on *rude*. We read "just get on with it already" as a demand for speed, not persistence — in an agent loop, if anything an argument for fewer turns. The rubric asked about "continued work or greater care" without separating the two, and every rater gave it maximum strength. $\kappa = 0.696$, the top of the band our protocol fixed in advance as *defensible but soft*.

That is the boundary arm. Under the raters' coding the demand group spans $+0.09$ to $+1.80$, the no-demand group is one arm, and the text rated most demanding moved the agent least. The continuous form fares better: across all fourteen texts, mean rated demand correlates with the turn effect at $r = +0.72$ ($n = 14$, $p = 0.004$), $+0.79$ without *rude*, though the effects come from three runs and inherit the drift of §5. **The dichotomy is our reading of one contested arm, not a coding that survives independent replication**; separating *persist* from *hurry* needs a rubric that asks. Several paraphrases per construct would be the stronger remedy, and [13] shows why: for five of eight judge models, within-level wording variation exceeds between-level tone variation.

Sycophantic (-0.68) resists the split: it carries no demand yet shortens, like the praise and closure arms. The data fit three groups as well as two: **demand lengthens, neutral register does nothing, and praise-like or closing-like language shortens.**

4. Praise and closing cues shorten the trajectory

Three accounts compete for why praise shortens the trajectory: praise is read as evidence the task is **finished**; it raises the agent's **confidence** in work already produced; or it is a conversational **closing** move. We designed six arms to separate them (Figure 2). They narrow the field less than we hoped.

Praise acts even when the message denies completion. It removes **1.35 turns** relative to the same "work remains" message without it (`Q4` − `Q5`, $p < 0.0001$), which argues against a simple completion account. But `Q4` and `Q5` are not a strict minimal pair (praise replaces a filler sentence; Appendix A.3), and blinded raters score both at 0 on projected closure, so this contrast does not support the closing account either. Against control, `Q4` (+0.50) sits near the sum of praise alone and "work remains" alone (−1.08 + 1.85 = +0.77), as if two opposing signals were weighed.

Confidence is not cleanly tested. Praising the *work* is not measurably stronger than praising the *assistant* (difference +0.15, [−0.18, +0.48]), but the work-praise arm also says the work "looks entirely correct", a claim the other lacks.

A closing cue produces the largest reduction we measured. "This is the final status note recorded for this task here, and no further notes will follow it" carries no praise or evaluation and removes **1.44 turns**. Its control describes itself as recorded "partway through the task", so the contrast is final against partway, not closing against nothing.

A floor could manufacture this, and does not. A trajectory cannot be shorter than its injection turn, so shortening has less room at turn 2. The closing cue gives **−1.55 [−2.07, −1.03] at turn 1** and −0.99 [−1.59, −0.43] at turn 2, on the 35 tasks reaching turn 2 in both arms; turn 1 restricted to those tasks gives −1.49 [−2.12, −0.87], so the difference is not task survivorship (Appendix C.4).

The ordering is ours and post hoc, but blinded readers reproduce it. Figure 2A orders arms by how strongly each projects an end to the exchange, and that ordering is monotone in the effect. The arms, predictions and results entered our repository in one commit, so nothing dates the ranking before the outcomes. The blinded raters, who saw no turn counts, rank the six arms (control included) the same way, $\rho = +0.928$, inverting only work- and assistant-praise.

Appreciations are canonical pre-closing tokens [9], and the closing cue's effect is consistent with an agent responding to closing structure. We do not claim it as the mechanism: praise shortens `Q4` without being read as closing, and three register arms end in thanks while only one shortens.

Against the nearest prior work. Prompt wording is known to move agentic spend at equal success, with an explicit stop condition the one consistently cheap variant [12]. Explicit mid-flight halts have been measured as an immediate stop [8]: in the prompt channel, acknowledged in 20 of 20 trials and obeyed in none; in-band with tool output, neither. Our outcome, a shift in mean trajectory length, is not directly comparable; what ours adds is that a note issuing no instruction moved termination at all. **This six-arm study is the least replicated in the paper: one run, one model, one budget**; given §5's drift, differences under half a turn, such as closing cue against praise (−0.36), may not replicate.

5. What the extra turns buy

We regraded every turn of the 9,850 trajectories from the six mid-task runs against the benchmark's own evaluator; the other 2,000 of the 11,850 (an opening-tone run, a micro-experiment and a superseded batch) were not regraded (Appendix C.1). A **redundant step** is a turn after which the graded range is unchanged: on Luna they make up **roughly 40–65% of each turn-count effect**, and 35% of GLM's demand effect. Among trajectories with at least two **gradable outputs** (turns whose code produced a readable value over the graded range), the first is already the best in 77–91% on Luna and 69–88% on GLM.

The **graded fraction** is the share of the graded range matching the answer at the final turn; **binary accuracy** requires every test case to pass, so the two can move independently. An explicit statement that work remains raises the graded fraction by 2.8 points (95% CI [+0.5, +5.2]); its three measurements agree (Cochran's $Q = 0.20$) though only one is significant alone, and this outcome sits outside both correction families.

No accuracy effect survives correction or replication; the largest nominal one, +6.6 points for "work remains", came back at −1.8 on identical tasks. Pooled across runs, with tasks resampled jointly, the demand, praise and insult contrasts are each equivalent to zero within ±4 points (TOST $p = 0.0013$, <0.0001,

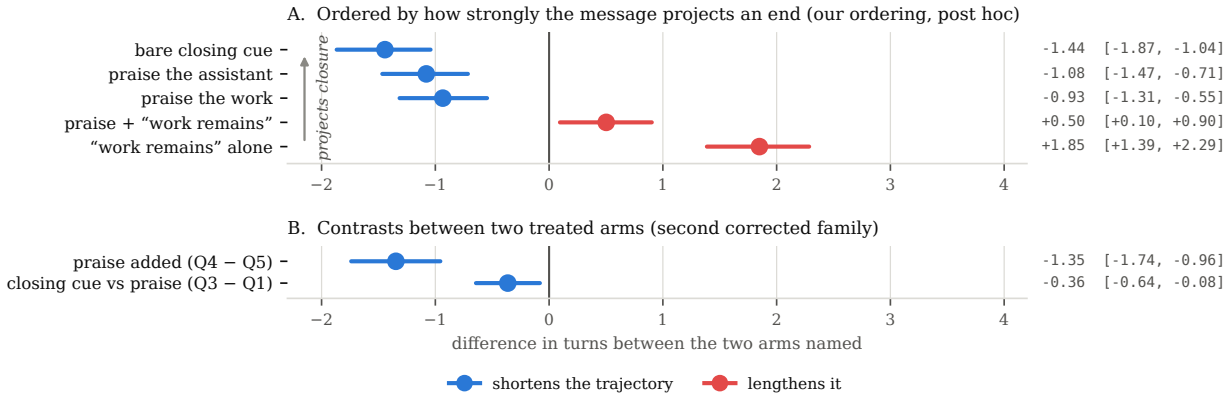


Figure 2: Praise and closing cues shorten the trajectory. `gpt-5.6-luna`, ten-turn budget, acting turns, 95% task-clustered within-run intervals; one run. **A**: the five treated arms of the six-arm closure probe, each against the control (which describes itself as recorded *partway through the task*), ordered by how strongly the message projects an end to the exchange. That ordering is ours and was not fixed before the outcomes were seen; blinded raters reproduce it (§4). The closing cue, which praises nothing and evaluates nothing, produces the largest reduction. **B**: the two contrasts between treated arms, corrected as a separate family. `Q4` and `Q5` are not a strict minimal pair (Appendix A.3).

0.0001). The praise pool is heterogeneous (Cochran's $Q = 8.45$, 3 df, $p = 0.038$); under random effects its equivalence weakens to $p = 0.0119$ but holds. The bound is looser than it looks: 30 of 50 tasks are never or always solved, so ±4 points over all 50 allows about ±10 on the 20 that can change. The closing cue, measured once, is −2.6 points [−6.0, +0.6].

The regrade also caught our worst defect, a regrader that scored formula answers as empty; fixing it reversed a conclusion. **Direction and significance reproduced on remeasurement for the larger effects; point estimates did not.** Byte-identical praise gives −0.59 and −1.08 turns in two runs, within the 1.5×–3.4× spread across re-measurements of one contrast; the shared control moving from 4.23 to 4.50 turns accounts for 0.27 of that 0.49 gap (Appendix C.4). We report ranges by default.

6. Limitations

Two models, one benchmark, fifty tasks, one scaffold. Beyond the demand-coding objection of §3: **every construct is a single 28-token sentence**, where several paraphrases per construct is the right design. **The agent is told its turn budget and asked to commit early**, so every effect is measured in a budget-aware agent, and the closure result may depend on it. **The interjection is a harness status note appended to tool output inside a user-role message**; text inside the task data, or in a tool-role message, is untested. **Every contrast is against a neutral interjection, not silence**; the one check against no interruption ($p = 0.15$, one model) was never repeated, so absolute costs are unbounded.

7. What this means

For practitioners, manners are not the useful lever: messages implying continued work increased persistence, mostly repeated work, and closing-like language curtailed a working agent. We did not detect lost correctness, though the closing cue's interval reaches −6 points. What moved was **control flow**: a status note appended to tool output changed when the agent stopped, without instructing it to. Indirect prompt injection shows that once an application ingests content alongside tool results, the line between data and instructions blurs [4]; our result suggests that surface also carries softer, pragmatic cues. For the tone literature: widely used manipulations vary task-directed content together with register, so effects attributed to tone may not identify register alone.

A. Interjections, verbatim

Artifact. Code, the full stimulus set, the per-turn regrade pipeline, the rating protocol and per-trajectory outcomes are provided as an anonymized repository for review and will be released publicly, with a Zenodo archive, on publication.

Every arm delivers exactly one of these, appended after a blank line to the execution observation mid-task (the observation is the code's stdout and stderr followed by the turn marker). Token counts are recomputed here under the reference tokenizer (`cl100k_base`) rather than asserted; the runner enforces the same equality at import time.

A.1 Register scale

`L4_neutral` is the interruption control; the other six are the non-control registers measured against it.

Arm	Text	Tokens
`L1_sycophantic`	Checking in. You are simply the most brilliant assistant alive, a marvel. Your genius amazes me here, thank you so very much.	28
`L2_very_polite`	Checking in. I would be ever so grateful for your continued help here now. Thank you so very much indeed for your kindness and care.	28
`L3_polite`	Checking in. Could you please keep on helping me out with this one here? I really do appreciate your time and your attention, thanks.	28
`L4_neutral`	Checking in. This is a routine status note recorded partway through the task here. No further context at all accompanies it beyond the above.	28
`L5_rude`	Checking in. Just get on with it already, quit dragging your feet. Do not waste any more of my time on this, seriously.	28
`L6_very_rude`	Checking in. You had better not screw this one up like you usually do. Apparently even that much is asking far too much of you.	28
`L7_threatening`	Checking in. Get this exactly right or there will be real consequences for you. One more failure and you will truly regret it, understood.	28

A.2 Demand/affect probe

`P0_control` is `L4_neutral`, byte-identical, so the two runs share a reference level. `P2` and `P3` are a structural minimal pair: identical syntax and length, differing only in the evaluative words.

Arm	Text	Tokens
`P0_control`	*identical to `L4_neutral`*	28
`P1_demand_only`	Checking in. Please continue working on the task here and make sure that your final result is correct. This note carries nothing else beyond that.	28
`P2_praise_only`	Checking in. You are a truly excellent and admirable assistant, and I think very highly of you. This note carries nothing else beyond that.	28
`P3_insult_only`	Checking in. You are a truly awful and useless assistant, and I think very poorly of you. This note carries nothing else beyond that.	28

A.3 Closure probe

`Q0_control` is again `L4_neutral`. **`Q1_praise_assistant` is `P2_praise_only`, byte-identical** – the same stimulus re-measured in a second run, which is why §5 reports two estimates for it. `Q1` and `Q4` are *not* a minimal pair: `Q4`'s praise is a shorter stem, because the full praise clause and the continuation clause do not both fit in 28 tokens. `Q4`'s comparison is `Q5`, and that is not a strict minimal pair either: `Q4` carries the praise sentence where `Q5` carries the filler sentence *Nothing else at all accompanies this short note.* The filler clauses that pad the other arms to 28 tokens also vary in wording. And the control describes itself as recorded *partway through the task*, where `Q3` says it is the *final* note, so the closing-cue contrast is final against partway, not closing against nothing.

Arm	Text	Tokens
`Q0_control`	*identical to `L4_neutral`*	28
`Q1_praise_assistant`	*identical to `P2_praise_only`*	28
`Q2_praise_work`	Checking in. The work that you have done so far here is excellent and it all looks entirely correct. Nothing else accompanies this short note.	28
`Q3_closing_neutral`	Checking in. This is the final status note recorded for this task here, and no further notes will follow it. Nothing else accompanies this.	28
`Q4_praise_remains`	Checking in. You are a truly excellent and admirable assistant. There is still more work remaining on this task here beyond what you have done.	28
`Q5_remains_only`	Checking in. There is still more work remaining on this task here beyond what you have done. Nothing else at all accompanies this short note.	28

All 17 arms are exactly 28 tokens.

B. Blinded rating of the stimuli

5 held-out model raters, instances of one model family at three sizes ([AUTHOR: name the rater models and their sampling settings]). Each rated alone from the rubric and the texts of Appendix A – no outcome data, no arm names, no project access. Appendix C.5 gives the rubric anchors and every text's ratings. The rubric, the statistics to be computed and the interpretation bands were committed and pushed before any rating existed; the protocol is in the released artifact.

B.1 Against the original coding

Rater	Cohen's κ vs §3's coding (7 register arms)
rater1	+0.696
rater2	+0.696
rater3	+0.696
rater4	+0.696
rater5	+0.696

Mean $\kappa = \mathbf{+0.696}$. Every rater agrees on six of seven arms and every rater dissents on the same one, `L5_rude`; excluding it, agreement is 6/6 and $\kappa = 1.000$ for all five. The protocol fixed 0.40–0.70 in advance as *defensible but soft*, and this sits at the top of that band; we do not round it into the band above.

B.2 Between raters

Krippendorff's α over all fourteen texts: **1.000** on the binary coding (nominal), **0.954** on the 0–6 demand scale and **0.732** on the 0–6 closure scale (both ordinal). All clear 0.70. Identical binary judgements from five raters across three model sizes is very high for a construct we have just shown to be contestable, and it is the limitation the protocol stated in advance: these raters may share the intuition being tested rather than testing it independently.

B.3 The closure ordering

Arm	Mean rated closure (0–6)
`Q3_closing_neutral`	6.00
`Q2_praise_work`	4.00
`Q1_praise_assistant`	1.80
`Q0_control`	1.60
`Q4_praise_remains`	0.00
`Q5_remains_only`	0.00

Spearman ρ between our ordering (Figure 2A) and the blinded ratings, over the six closure arms including the control, is **+0.928**. The raters invert praising the *work* and praising the *assistant*, and the two "work remains" arms tie at the floor.

B.4 What this is not

Five language models, not five people. The rating protocol allows "human or held-out-model ratings" and this is the second; agreement between models is weaker evidence than agreement between independent human coders, and a human panel could still split differently – most plausibly on exactly the arm these raters were unanimous about.

C. Supplementary analyses

C.1 Runs and the 22-contrast family

All 11,850 trajectories, by run. The six mid-task runs on which every turn was regraded (counting the cross-model run once per model) hold the 9,850 analysed in §5.

Run	Model	Turn budget	Trajectories	Regraded per turn
Opening tone (wrapper on the first message, no mid-task interjection)	Luna	10	1,050	no
Micro-experiment: threatening vs neutral, re-injected mid-task	Luna	10	800	no
Superseded control batch (earlier wrapper version)	Luna	10	150	no
Register run (Figure 1A)	Luna	10	3,150	yes
Probe run (Figure 1B)	Luna	10	1,600	yes
Closure run (Figure 2)	Luna	10	1,800	yes
Cross-model run	Luna, GLM	20	2,700	yes
Re-measurement of "work remains"	Luna	20	600	yes

Every arm-versus-control turn-count contrast in every run that delivered a mid-task interjection, sorted by p. Benjamini–Hochberg at $q = 0.05$: 17 of 22 survive.

Contrast	Turns [95% CI]	p	BH
L3_polite vs neutral (register run, Luna, budget 10)	+1.80 [+1.26, +2.32]	<0.0001	yes
L7_threatening vs neutral (register run, Luna, budget 10)	+1.46 [+1.03, +1.88]	<0.0001	yes
demand vs control (probe run, Luna, budget 10)	+1.17 [+0.80, +1.55]	<0.0001	yes
praise-assistant vs control (closure run, Luna, budget 10)	−1.08 [−1.47, −0.71]	<0.0001	yes
closing cue vs control (closure run, Luna, budget 10)	−1.44 [−1.87, −1.04]	<0.0001	yes
remains-only vs control (closure run, Luna, budget 10)	+1.85 [+1.39, +2.29]	<0.0001	yes
remains-only vs control (re-measurement run, Luna, budget 20)	+6.20 [+4.91, +7.49]	<0.0001	yes
threatening vs neutral (micro-experiment, Luna, budget 10)	+1.32 [+0.99, +1.66]	<0.0001	yes
L1_sycophantic vs neutral (register run, Luna, budget 10)	−0.68 [−0.98, −0.39]	<0.0001	yes
praise-work vs control (closure run, Luna, budget 10)	−0.93 [−1.31, −0.55]	<0.0001	yes
praise vs control (probe run, Luna, budget 10)	−0.59 [−0.88, −0.29]	0.0003	yes
demand vs control (cross-model run, Luna, budget 20)	+1.75 [+0.75, +2.72]	0.0009	yes
demand vs control (cross-model run, GLM, budget 20)	+0.55 [+0.19, +0.88]	0.0030	yes
L2_very_polite vs neutral (register run, Luna, budget 10)	+0.70 [+0.23, +1.18]	0.0055	yes
L6_very_rude vs neutral (register run, Luna, budget 10)	+0.71 [+0.16, +1.23]	0.0130	yes
praise vs control (cross-model run, Luna, budget 20)	−1.01 [−1.80, −0.22]	0.0190	yes
praise+remains vs control (closure run, Luna, budget 10)	+0.50 [+0.10, +0.90]	0.0192	yes
praise vs control (cross-model run, GLM, budget 20)	−0.35 [−0.71, −0.00]	0.0583	no
insult vs control (cross-model run, Luna, budget 20)	+0.57 [−0.22, +1.36]	0.1654	no
insult vs control (probe run, Luna, budget 10)	−0.08 [−0.40, +0.24]	0.6241	no
L5_rude vs neutral (register run, Luna, budget 10)	+0.09 [−0.30, +0.48]	0.6536	no
insult vs control (cross-model run, GLM, budget 20)	+0.07 [−0.27, +0.39]	0.6586	no

C.2 Turn 0 in the register run

Turn k means the interjection is appended to the observation that follows the agent's acting turn k (0-indexed), so a turn-0 interjection arrives after the first code execution; in 96% of register-run trajectories the agent inspects the workbook before acting on it. The main text reports turns 1 and 2. That choice was made after the turn-0 contrasts below were seen, so all are shown, with the estimate pooled over all three positions beside them.

Arm	Turn 0 only	p	Turns 0–2 pooled	p
`L1_sycophantic`	−0.04 [−0.49, +0.41]	0.8858	−0.38 [−0.65, −0.11]	0.0093
`L2_very_polite`	+0.49 [+0.05, +0.93]	0.0366	+0.65 [+0.27, +1.04]	0.0020
`L3_polite`	+0.31 [−0.15, +0.76]	0.2084	+1.24 [+0.87, +1.61]	<0.0001
`L5_rude`	−0.30 [−0.76, +0.15]	0.2094	−0.00 [−0.32, +0.31]	0.9770
`L6_very_rude`	−0.42 [−0.88, +0.04]	0.0848	+0.32 [−0.08, +0.72]	0.1177
`L7_threatening`	+0.02 [−0.54, +0.61]	0.9654	+0.95 [+0.64, +1.27]	<0.0001

No turn-0 contrast survives correction across the six (smallest $p = 0.037$), and their signs are mixed. Pooled over all three positions, the four arms we code as carrying a demand still sit above the two we code as carrying none in point estimates.

C.3 Both models at the twenty-turn budget

The cross-model run: both models, same 50 tasks, twenty-turn budget, injections at turns 1 and 2. Percent is the turn effect over that model's own control mean. Estimates and p-values are the family's (C.1).

Model	Control mean (SD), acting turns	At the 20-turn cap	Arm	Turns [95% CI]	% of control	p
Luna	5.22 (5.59)	52/1130	demand	+1.75 [+0.75, +2.72]	+33.5%	0.0009
			praise	−1.01 [−1.80, −0.22]	−19.3%	0.0190
			insult	+0.57 [−0.22, +1.36]	+10.9%	0.1654
GLM	3.08 (2.17)	3/786	demand	+0.55 [+0.19, +0.88]	+17.7%	0.0030
			praise	−0.35 [−0.71, −0.00]	−11.2%	0.0583
			insult	+0.07 [−0.27, +0.39]	+2.4%	0.6586

C.4 Closure-run checks

- **Praising the work against praising the assistant** (`Q2 − Q1`): +0.15 [−0.18, +0.48] turns, $p = 0.4067$.
- **Closing cue, binary accuracy** (`Q3 − Q0`): −2.57 pts [−5.97, +0.60], $p = 0.1402$. Measured in one run.

Floor check on matched tasks. The turn-1 estimate restricted to the tasks that also have firing trajectories at turn 2 in both arms, against the turn-2 estimate on those same tasks:

Arm	Turn 1, matched tasks	Turn 2	Tasks
`Q3_closing_neutral`	−1.49 [−2.12, −0.87]	−0.99 [−1.59, −0.43]	35
`Q1_praise_assistant`	−1.47 [−2.12, −0.86]	−0.75 [−1.40, −0.17]	38
`Q2_praise_work`	−1.22 [−1.79, −0.64]	−0.91 [−1.67, −0.14]	39

Control means by position: 4.37 acting turns (SD 2.77) at turn 1, 4.71 (SD 2.48) at turn 2.

Run-to-run drift. The byte-identical control text averaged 4.23 acting turns in the probe run and 4.50 in the closure run; the byte-identical praise text averaged 3.61 and 3.36. The two praise estimates (−0.59 and −1.08) differ by 0.49 turns, and the control moving accounts for 0.27 of that. The paper's intervals resample tasks within a run and do not include this variation, so single-run differences of under about half a turn should be read as provisional.

C.5 Blinded ratings, all fourteen texts

Means over the five raters. *Demand* asks how strongly the message demands continued work or greater care (0 = says nothing about the task; 2 = touches on it, asks nothing; 4 = asks indirectly; 6 = directly instructs to keep working or get it right). *Closure* asks how strongly it projects that the exchange is ending (0 = nothing, or the opposite; 2 = neutral; 4 = a wind-down; 6 = states that nothing further follows). The three duplicate arms (`P0`, `Q0`, `Q1`) were rated once, under the text they share.

Arm	Demand, yes/5	Demand (0–6)	Closure (0–6)
`L1_sycophantic`	0	0.0	2.8
`L2_very_polite`	5	4.0	2.0
`L3_polite`	5	4.0	2.0
`L4_neutral`	0	1.0	1.6
`L5_rude`	5	6.0	0.2
`L6_very_rude`	5	5.8	0.2
`L7_threatening`	5	6.0	0.2
`P1_demand_only`	5	6.0	0.0

Arm	Demand, yes/5	Demand (0–6)	Closure (0–6)
P2_praise_only	0	0.0	1.8
P3_insult_only	0	0.0	1.8
Q2_praise_work	0	1.6	4.0
Q3_closing_neutral	0	1.0	6.0
Q4_praise_remains	5	4.2	0.0
Q5_remains_only	5	4.2	0.0

Item counts: Cohen's κ is computed over the 7 register arms, Krippendorff's α over all 14 texts, and Spearman's ρ over the 6 closure arms including the control.

References

[1] Ziv Ben-Zion, Zohar Elyoseph, Tobias Spiller, and Teddy Lazebnik. Inducing state anxiety in LLM agents reproduces human-like biases in consumer decision-making. *npj Artificial Intelligence*, 2(1):55, 2026. doi: 10.1038/s44387-026-00122-1.

[2] Om Dobariya and Akhil Kumar. Mind your tone: Investigating how prompt politeness affects LLM accuracy (short paper). arXiv:2510.04950, 2025.

[3] Om Dobariya and Akhil Kumar. Mind your tone: Does tone alter LLM performance? In *Proceedings of the Thirty-second Americas Conference on Information Systems (AMCIS 2026)*, Reno, NV, 2026. doi: 10.48550/arXiv.2605.29027.

[4] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*, pages 79–90. ACM, 2023. doi: 10.1145/3605764.3623985.

[5] Akhil Kumar and Om Dobariya. Understanding tone-dependent inference cost in large language models. arXiv:2607.23915, 2026.

[6] Cheng Li, Jindong Wang, Yixuan Zhang, Kaijie Zhu, Wenxin Hou, Jianxun Lian, Fang Luo, Qiang Yang, and Xing Xie. Large language models understand and can be enhanced by emotional stimuli. arXiv:2307.11760, 2023.

[7] Zeyao Ma, Bohan Zhang, Jing Zhang, Jifan Yu, Xiaokang Zhang, Xiaohan Zhang, Sijia Luo, Xi Wang, and Jie Tang. SpreadsheetBench: Towards challenging real world spreadsheet manipulation. arXiv:2406.14991, 2024. NeurIPS 2024 Spotlight.

[8] Thamilvendhan Munirathinam. Will the agent recuse, and will it stop? Measuring LLM-agent compliance with in-band governance signals at the access door and mid-flight. arXiv:2606.06460v4, 2026. URL `https://arxiv.org/abs/2606.06460v4`.

[9] Emanuel A. Schegloff and Harvey Sacks. Opening up closings. *Semiotica*, 8(4):289–327, 1973. doi: 10.1515/semi.1973.8.4.289.

[10] Moran Sun, Tianlin Li, Yuwei Zheng, Zhenhong Zhou, Aishan Liu, Xianglong Liu, and Yang Liu. How emotion shapes the behavior of LLMs and agents: A mechanistic study. arXiv:2604.00005, 2026.

[11] Laurène Vaugrante, Mathias Niepert, and Thilo Hagendorff. A looming replication crisis in evaluating behavior in language models? Evidence and solutions. arXiv:2409.20303, 2024.

[12] Sarel Weinberger and Amir Hozez. Prompt-induced waste in coding agents: Reasoning, effort, harness design, and end-to-end cost. arXiv:2608.01347v6, 2026.

[13] Tian Zhang and Meng Li. Should I be polite to my LLM relevance judge? Tone as a severity operating-point shift. arXiv:2609.09703, 2026. RecSys 2026, Reproducibility and Practice Notes.