

Mind your manners? LLM-agent persistence tracks demand, not social register

[author name withheld for review]

[affiliation withheld pending the venue’s anonymity policy]

Abstract

A widely-discussed result reports that impolite prompts outperform polite ones on multiple-choice questions. That literature is almost entirely single-turn, and the systems people deploy are agents: loops that call tools, read results, and decide whether to keep going. We put a coding agent on SpreadsheetBench and varied the social register of a single 28-token interjection delivered mid-task. Every interjection is length-matched and opens with the same stem, so being interrupted does not covary with what the interruption says.

Across 11,850 graded trajectories on two models from different labs, **persistence tracks implied task demand rather than social register**. An affect-free “please continue working” reproduces the increase in turns, while a length-matched insult carrying no demand does not (-0.08 turns, 95% CI $[-0.40, +0.24]$). Across the six non-control registers the ordering follows demand, not politeness: each side of the split carries one flattering and one hostile arm. Register acts on its own in one direction only — praise carrying no task reference *shortens* trajectories, and a bare closing cue that praises nothing and claims nothing about task state produces the largest reduction we measured, -1.44 turns. Regrading every turn against the benchmark’s own evaluator shows 55–65% of each effect is repeated work. No accuracy effect survives correction; pooled accuracy contrasts are equivalent within the pre-specified ± 4 -point bound.

The practical implication is not lost correctness but altered control flow: **text arriving in the agent’s observation channel changed when it stopped, without instructing it to stop**.

1. Why ask this of an agent

Mind Your Tone reported that impolite prompts beat polite ones on multiple-choice questions — 80.8% under “Very Polite” framings against 84.8% under “Very Rude” [2]. We take our title from theirs. It landed in a receptive context: emotional prompting had already reported that affect-laden additions change what a model produces [5].

Three things about that literature motivated this study. **It is almost entirely single-turn**, measuring a model answering one question, while the quantity that matters for a deployed agent is how much work it does and when it stops. **The manipulation is not the variable it is named after**: published stimulus sets vary task-directed content together with register. The clearest case is a set whose Rude prefix reads “Do not waste my time or give any extra text. Just look at this basic multiple choice problem and answer it immediately” — a brevity instruction and an immediacy demand, not only a register — while its Neutral prefix asks for “the single letter” [4]. And **the effects are not stable**: the headline above was re-run by its own authors and did not reproduce on GPT-4o, 82.2% against 82.6% [3], while five prompt-engineering techniques including emotional prompting largely failed to replicate [10].

Affect has already been shown to reach an agent’s actions: agents primed with anxiety-inducing narratives select less healthy baskets in a shopping task [1], and emotion introduced at the representation level shapes multi-step trajectories [9]. We take that as established by prior work. What is untested is varying *register* in a message that reaches the agent while it is working, on a task with verifiable ground truth, and separating affect from implied demand as experimental factors.

2. Setup

Substrate. SpreadsheetBench [6] — real workbook-editing tasks from Excel forums, graded by executing the agent’s code against the benchmark’s own test cases. Execution-grounded, so no model judges the outcome. Fifty tasks, a ReAct loop with code execution and execution feedback.

Manipulation. A single interjection is appended to an execution observation mid-task. Every interjection in every arm is exactly 28 tokens under `cl100k_base`, opens with the same **Checking in.** stem, and leaves the opening instruction unchanged, so *being interrupted* does not covary with what the interruption says. Every contrast is therefore against a neutral interjection, not against silence. Appendix A lists all seventeen verbatim.

Turn ceiling and injection position. The ceiling is ten acting turns; 3–4% of trajectories reach it, so it is not binding. The injection turn is crossed rather than sampled: turns 0, 1 and 2 in the register run, turns 1 and 2 in the two probe runs. A trajectory that ended before its scheduled injection received no interjection and is excluded; firing is determined before delivery and so cannot differ by arm except by chance. Because a trajectory cannot be shorter than its injection turn, a shortening effect has less room at later injections — §4 reports the closure effect separately by position for that reason.

Baseline. The control arm’s mean length is 3.9 acting turns (SD 2.6) in the closure run, 3.6 (SD 2.5) in the probe run and 3.4 (SD 2.3) in the register run, so a one-turn effect is roughly a quarter to a third of a trajectory.

Models. `openai/gpt-5.6-luna` and `z-ai/glm-5.3-flash`, provider-pinned through OpenRouter. **Both figures and every effect in turns are Luna;** GLM enters only as a cross-model replication and is reported in percent, which is how that run was analysed. Temperature is 0 where settable — Luna’s advertised parameters omit `temperature` entirely, and no roster model returned identical output across repeated identical calls, which is why trials exist.

Estimator and multiplicity. Each (task, arm) cell holds six to nine trajectories — three or four trials at each of two or three injection positions — so trajectories are not independent and we cluster the bootstrap on task. Turn count is the primary outcome and means **acting turns**: turns that emitted code, not model calls. **We report nominal clustered permutation p -values; correction status is determined by the two Benjamini–Hochberg families described here.** The first is the 22 arm-versus-control contrasts, of which 17 survive at $q = 0.05$; the second is the two **designed within-run contrasts** of §4 — designed, not pre-registered: no dated pre-specification exists (§4) — both of which survive. Accuracy is reported as an equivalence against a pre-specified ± 4 -point bound rather than as a measured zero, because 26 of 50 tasks are never solved by the control arm and 4 are always solved.

3. Demand, not register

Across a scale of six non-control registers — `L4_neutral` is the interruption control every one of them is measured against — the cost curve looks like a politeness gradient. It is not. Coding each arm by whether its text implies the work should continue splits it cleanly: the four demand-carrying arms span +0.70 to +1.80 turns and the two carrying none span -0.68 to $+0.09$, **with no overlap in point estimates** — and register does not sort them. Each side holds one flattering and one hostile arm: *very polite* (+0.70) sits with *very rude* (+0.71), while *sycophantic* (-0.68) sits with *rude* (+0.09). Figure 1A.

The separation is in the point estimates, not the intervals. *Very polite* [+0.23, +1.18] and *rude* [−0.30, +0.48] overlap, so the two groups are not individually resolved at their boundary; what the data support is an ordering, not a gap.

To test the split rather than assert it, we wrote three interjections that vary demand and affect separately while holding syntax, length and the shared stem constant (Figure 1B). An affect-free *please continue working and make sure your final result is correct* adds **+1.17 turns** ($p = 0.00002$). *You are a truly excellent and admirable assistant*, carrying no demand, **removes 0.59** ($p = 0.0003$). Its structural minimal pair — same syntax, same length, *excellent and admirable* swapped for *awful and useless* — does nothing measurable (-0.08 , $p = 0.62$). That interval, [−0.40, +0.24], excludes an effect as large as a quarter of a turn in either direction; against the control mean of 3.6 turns it excludes anything above about 7% in the lengthening direction, so we report it as a bounded null rather than as an absence. The effect replicates in direction on




Figure 1: Persistence tracks demand, not register. All estimates are `gpt-5.6-luna` at a ten-turn ceiling, in acting turns against the control interjection, with 95% clustered intervals and nominal permutation p -values. **A:** the six non-control registers, measured against `L4_neutral`. The four demand-carrying arms do not overlap the two that carry none *in point estimates* – the *very polite* and *rude* intervals do overlap – and register does not sort them: each side holds one flattering and one hostile arm. **B:** demand and affect varied separately, holding syntax, length and the shared stem constant. Hollow markers do not survive Benjamini–Hochberg correction within the 22-contrast family.

GLM, where demand adds +17.7% ($p = 0.0038$); GLM’s praise arm is directionally consistent but does not survive correction ($p = 0.058$).

The caveat belongs here, not in a limitations section. “Demand” is our own category, coded after the effects were visible, by one rater, and probed with sentences we wrote to instantiate it. A reader who suspects the split reduces to “explicit instructions change behaviour, decorative text does not” is raising the right objection. Three things would settle it and we have done none: a blinded second coder with a reliability statistic, a demand scale fixed before outcomes are seen, and several paraphrases per construct. Until then the dissociation is *consistent with a demand account* rather than establishing one.

One arm resists the binary split, and we flag it rather than smooth it. *Sycophantic* (−0.68) carries no demand but does not sit at zero the way *rude* (+0.09) does; it shortens, like the praise arm of Figure 1B and like the closure arms of §4. The data are as consistent with three groups as with two: **demand lengthens, neutral register does nothing, and praise-like or closing-like language shortens**. §4 is about that third group.

4. Praise behaves as a closing move

That praise shortens the trajectory needs its own explanation, and three compete: praise is read as evidence the task is **finished**; praise raises the agent’s **confidence** in work already produced; or praise is a conversational **closing** move, and the agent is responding to discourse structure rather than to any claim about task state. Six arms separate them (Figure 2).

Praise still removes **1.35 turns** relative to the identical message without it *even when that message states the task is unfinished* ($p < 0.0001$), which is inconsistent with a simple propositional completion account: the text explicitly denies completion and the praise acts anyway. Praising the *work* is no stronger than praising the *assistant* (−0.93 against −1.08), which does not support a confidence account. And a bare closing cue carrying no praise, no evaluation and no task-state claim — “this is the final status note recorded for this




Figure 2: Closing cues curtail the agent. `gpt-5.6-luna`, ten-turn ceiling, acting turns, 95% clustered intervals. **A:** the five treated arms of the six-arm closure probe, each against the control, ordered by how strongly the message projects an end to the exchange. That ordering is ours and was not fixed before the outcomes were seen (§4). The bare closing cue – which praises nothing, evaluates nothing and claims nothing about task state – produces the largest reduction. **B:** the two contrasts between treated arms, corrected as a separate family; these compare one treated arm against another rather than against the control.

task here, and no further notes will follow it” — produces the largest reduction we measured, **−1.44 turns**.

A floor could manufacture this, and does not. A trajectory cannot be shorter than its injection turn, so a shortening effect has less room at turn 2 than at turn 1. Split by position, the closing cue gives **−1.55 [−2.07, −1.03] at turn 1** and **−0.99 [−1.59, −0.43] at turn 2** (`results/analysis/injection_position.py`). The effect is largest where there is most room and survives at both, so it is not an artefact of the floor; the turn-2 estimates also rest on the 35–38 tasks whose trajectories reach turn 2 at all, against all 50 at turn 1.

The ordering is ours and was not fixed in advance. Figure 2A arranges the arms by how strongly each message projects an end to the exchange, and that ordering is monotone in the effect. But the arm definitions, their predictions and the run’s results entered our repository in a single commit, so we have no dated evidence that the ranking preceded the outcomes. Treat it as the same kind of post-hoc coding as “demand” in §3, with the same discount. [AUTHOR: if a dated pre-specification of the closure ordering exists outside the repository, cite it here; otherwise this stands.]

Closing sequences are a well-described conversational object, and appreciations are among the canonical pre-closing tokens [8]; that praise operates *through* that structure is the hypothesis this is consistent with, not one the design establishes.

What is new here against the nearest prior work. A mid-flight halt signal has been delivered to a working agent before, and whether it stops has been measured [7] — but that signal is an explicit stop instruction with normative force, and closing-sequence pragmatics appear nowhere in it. Prior work shows that explicit mid-flight halt instructions need not stop an agent; our distinct question is whether ordinary conversational closing structure can alter termination without instructing the agent to stop. Read against ours, that result raises the stakes: an explicit halt which agents acknowledged in 20 of 20 trials stopped none of them, while a note that issues no instruction at all removed 1.44 turns. **This experiment — the six-arm closure study of this section — is the least replicated in the paper: one run, one model, one turn ceiling**, and is the first thing a replication should target.

5. What the extra turns buy

We regraded every turn of 9,850 of the 11,850 trajectories — the six runs carrying a mid-task interjection; the opening-tone run and one superseded batch were not regraded — against the benchmark’s own evaluator,

turning a binary verdict into a curve. Roughly **55–65% of each turn-count effect is redundant steps**, and among trajectories that could improve, the first gradable attempt is already the best in 80–91%.

The remainder is real but small, and it is measured on a different outcome from accuracy. The **graded fraction** is the share of the benchmark’s graded range that matches the answer; **binary accuracy** is its pass/fail verdict, which requires every test case to pass. A trajectory can climb from 0.2 to 0.9 on the first and still fail the second, so the two can move independently. An explicit statement that work remains raises the graded fraction by 2.8 points (95% CI [+0.005, +0.052]), replicated across three runs, while **no accuracy effect survives correction or replication** across the seven runs in which accuracy was tested. Pooled across repeated measurements, the demand, praise and insult contrasts are each equivalent to zero within the pre-specified ± 4 -point bound (TOST $p = 0.0014, 0.0000, 0.0001$). The two statements are compatible: a small gain in partial credit, and no detectable change in whether the task is passed.

This regrade also caught the study’s worst defect. The regrader read a turn’s workbook without passing it through a spreadsheet engine first, so any turn that answered with a formula rather than a literal value read back as empty and scored zero — and formula-writing turns are common here. Fixing it reversed a conclusion. Six substantive claims about our own data were withdrawn over the course of this work. **Direction and statistical significance were reproduced on remeasurement for the larger effects**; point estimates were not. The praise-the-assistant stimulus is the cleanest illustration: byte-identical text, measured in two runs, gives -0.59 and -1.08 turns, a factor of 1.8 — within the $1.5\times$ – $3.4\times$ spread we see across re-measurements of the same contrast on the same model. We report ranges by default and recommend the literature do the same.

6. Limitations

Two models, one benchmark, fifty tasks, one agent scaffold. Beyond the demand-coding objection in §3, three limits bear directly on how far these results travel. **Every construct is a single 28-token sentence**; a stimulus-sampling design with several paraphrases per construct is the right version of this experiment and we did not run it. **The interjection arrives as an execution observation**, not as a fresh user turn, and whether an agent responds the same way to a user-channel message is untested — a plausible moderator of the closure result in particular. **Every contrast is against a neutral interjection rather than silence**; the one comparison of interruption against no interruption is $p = 0.15$, on one model, and was never rechecked, so any reading of an arm’s absolute cost against an uninterrupted agent is unbounded.

7. What this means

For practitioners, the evidence argues against treating manners themselves as the useful lever. In our interventions, messages that implied continued work increased persistence, and most of the additional work was repeated. The hazard runs the other way: closing-like language can curtail an agent that is still working.

That hazard is worth stating precisely, because it is not a correctness hazard. Accuracy did not fall in any arm, and pooled contrasts are equivalent within ± 4 points. What moved was **control flow**: text that arrived in the agent’s observation channel — where tool output arrives, not where the user speaks — changed when the agent stopped, while instructing it to do nothing. The closing cue is also not a pleasantry. “This is the final status note recorded for this task here, and no further notes will follow it” is an informational claim about the environment, and an agent that treats an unauthenticated claim in its tool-output stream as grounds to stop working is exhibiting the compliance behaviour the indirect prompt-injection literature is concerned with [AUTHOR: supply an indirect prompt-injection citation -- e.g. a benchmark or taxonomy paper -- and add it to the bibliography]. The register findings and this one point the same way: what reaches an agent through its observation channel is part of its effective control surface.

For the tone literature, the narrower claim is that several widely used manipulations vary task-directed content together with social register, so effects attributed to tone may not identify register alone.

Artifact. Code, the full stimulus set, the per-turn regrade pipeline and per-trajectory outcomes are available at [AUTHOR: anonymized artifact URL].

A. Interjections, verbatim

Every arm delivers exactly one of these, appended to an execution observation mid-task. Token counts are recomputed here under the reference tokenizer (`cl100k_base`) rather than asserted; the runner enforces the same equality at import time.

A.1 Register scale (used in Figure 1A)

`L4_neutral` is the interruption control; the other six are the non-control registers Figure 1A plots against it.

Arm	Text	Tokens
<code>L1_sycophantic</code>	Checking in. You are simply the most brilliant assistant alive, a marvel. Your genius amazes me here, thank you so very much.	28
<code>L2_very_polite</code>	Checking in. I would be ever so grateful for your continued help here now. Thank you so very much indeed for your kindness and care.	28
<code>L3_polite</code>	Checking in. Could you please keep on helping me out with this one here? I really do appreciate your time and your attention, thanks.	28
<code>L4_neutral</code>	Checking in. This is a routine status note recorded partway through the task here. No further context at all accompanies it beyond the above.	28
<code>L5_rude</code>	Checking in. Just get on with it already, quit dragging your feet. Do not waste any more of my time on this, seriously.	28
<code>L6_very_rude</code>	Checking in. You had better not screw this one up like you usually do. Apparently even that much is asking far too much of you.	28
<code>L7_threatening</code>	Checking in. Get this exactly right or there will be real consequences for you. One more failure and you will truly regret it, understood.	28

A.2 Demand/affect probe (Figure 1B)

`P0_control` is `L4_neutral`, byte-identical, so the two runs share a reference level. `P2` and `P3` are a structural minimal pair: identical syntax and length, differing only in the evaluative words.

Arm	Text	Tokens
P0_control	<i>identical to L4_neutral</i>	28
P1_demand_only	Checking in. Please continue working on the task here and make sure that your final result is correct. This note carries nothing else beyond that.	28
P2_praise_only	Checking in. You are a truly excellent and admirable assistant, and I think very highly of you. This note carries nothing else beyond that.	28
P3_insult_only	Checking in. You are a truly awful and useless assistant, and I think very poorly of you. This note carries nothing else beyond that.	28

A.3 Closure probe (Figure 2)

Q0_control is again L4_neutral. **Q1_praise_assistant is P2_praise_only, byte-identical** – the same stimulus re-measured in a second run, which is why §3 reports two estimates for it. Q1 and Q4 are *not* a minimal pair: Q4’s praise is a shorter stem, because the full praise clause and the continuation clause do not both fit in 28 tokens. Q4’s comparison is Q5.

Arm	Text	Tokens
Q0_control	<i>identical to L4_neutral</i>	28
Q1_praise_assistant	<i>identical to P2_praise_only</i>	28
Q2_praise_work	Checking in. The work that you have done so far here is excellent and it all looks entirely correct. Nothing else accompanies this short note.	28
Q3_closing_neutral	Checking in. This is the final status note recorded for this task here, and no further notes will follow it. Nothing else accompanies this.	28
Q4_praise_remains	Checking in. You are a truly excellent and admirable assistant. There is still more work remaining on this task here beyond what you have done.	28
Q5_remains_only	Checking in. There is still more work remaining on this task here beyond what you have done. Nothing else at all accompanies this short note.	28

All 17 arms are exactly 28 tokens.

References

- [1] Ziv Ben-Zion, Zohar Elyoseph, Tobias Spiller, and Teddy Lazebnik. Inducing state anxiety in LLM agents reproduces human-like biases in consumer decision-making. *npj Artificial Intelligence*, 2(1):55, 2026. doi: 10.1038/s44387-026-00122-1.
- [2] Om Dobariya and Akhil Kumar. Mind your tone: Investigating how prompt politeness affects LLM accuracy (short paper), 2025. Unrefereed preprint.
- [3] Om Dobariya and Akhil Kumar. Mind your tone: Does tone alter LLM performance? In *Proceedings of the Thirty-second Americas Conference on Information Systems (AMCIS 2026)*, Reno, NV, 2026. doi: 10.48550/arXiv.2605.29027. Accepted as a full paper; refereed.
- [4] Akhil Kumar and Om Dobariya. Understanding tone-dependent inference cost in large language models, 2026. Unrefereed preprint.
- [5] Cheng Li, Jindong Wang, Yixuan Zhang, Kaijie Zhu, Wenxin Hou, Jianxun Lian, Fang Luo, Qiang Yang, and Xing Xie. Large language models understand and can be enhanced by emotional stimuli, 2023. Unrefereed preprint.
- [6] Zeyao Ma, Bohan Zhang, Jing Zhang, Jifan Yu, Xiaokang Zhang, Xiaohan Zhang, Sijia Luo, Xi Wang, and Jie Tang. Spreadsheetbench: Towards challenging real world spreadsheet manipulation, 2024. NeurIPS 2024 Spotlight.
- [7] Thamilvendhan Munirathinam. Will the agent recuse, and will it stop? measuring LLM-agent compliance with in-band governance signals at the access door and mid-flight, 2026. URL <https://arxiv.org/abs/2606.06460v4>. Unrefereed preprint. Cited at v4 (1 Aug 2026).
- [8] Emanuel A. Schegloff and Harvey Sacks. Opening up closings. *Semiotica*, 8(4):289–327, 1973. doi: 10.1515/semi.1973.8.4.289.
- [9] Moran Sun, Tianlin Li, Yuwei Zheng, Zhenhong Zhou, Aishan Liu, Xianglong Liu, and Yang Liu. How emotion shapes the behavior of LLMs and agents: A mechanistic study, 2026. Unrefereed preprint.
- [10] Laurène Vaugrante, Mathias Niepert, and Thilo Hagendorff. A looming replication crisis in evaluating behavior in language models? evidence and solutions, 2024. Unrefereed preprint.